

LAB 6 - RCNN AND FASTER RCNN

Dr. Pandiyaraju V
Shravan Venkatraman

September 27, 2024

TABLE OF CONTENTS

01

INTRODUCTION

02

OBJECT DETECTION

03

R-CNN

04

FAST R-CNN

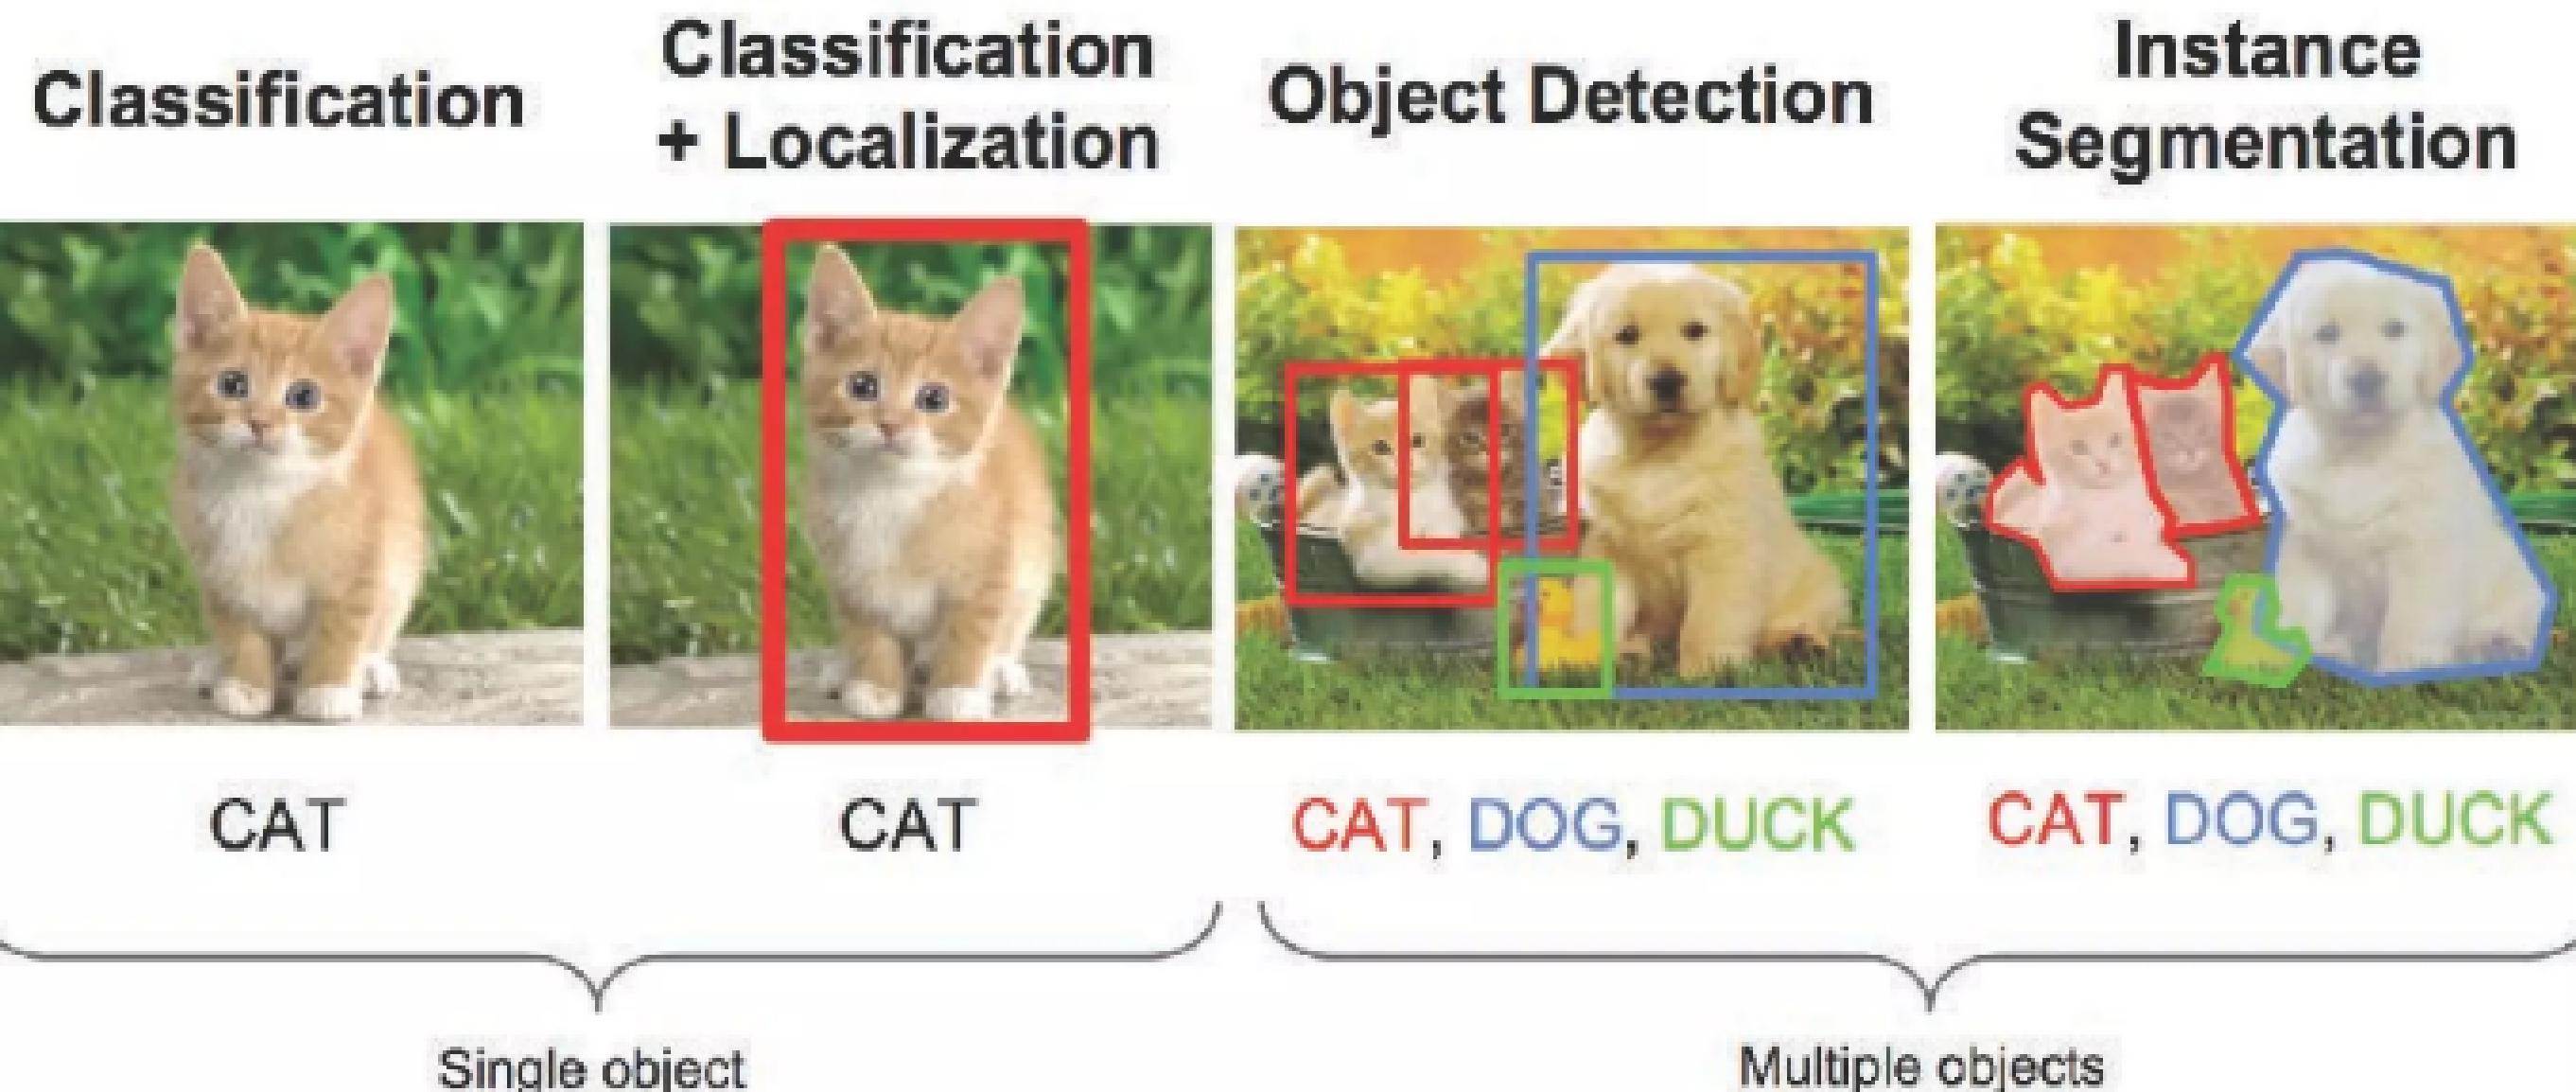
05

FASTER R-CNN


06

YOLO


COMPUTER VISION TASKS



R-CNN: REGIONS WITH CNN FEATURES



1. Input
image



2. Extract region
proposals (~2k)

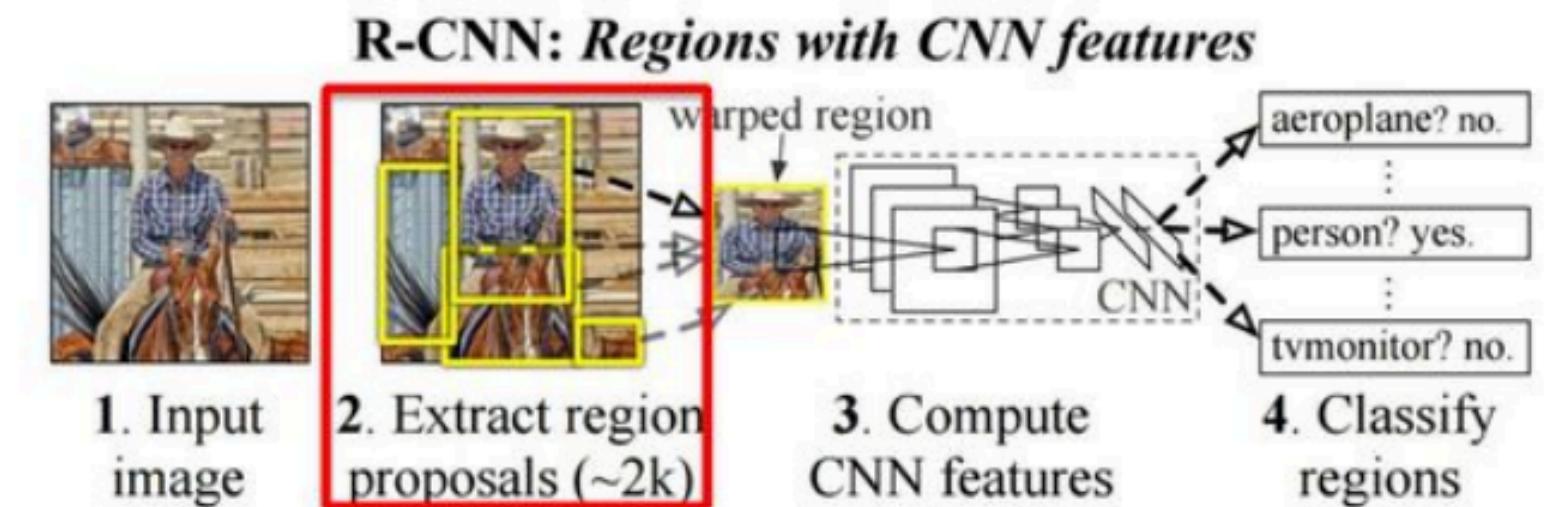
3. Compute
CNN features

4. Classify
regions


aeroplane? no.
⋮
person? yes.
⋮
tvmonitor? no.

R-CNN: REGIONS WITH CNN FEATURES

- First Stage: generate category-independent region proposals.
 - 2000 Region proposals every image.





Selective Search: Combining strengths of both exhaustive search and segmentation.



R-CNN: REGIONS WITH CNN FEATURES

- First Stage: generate category-independent region proposals.
 - 2000 Region proposals every image.
- Second Stage: Extract a fixed length feature vector from each region.
 - a 4096-dimensional feature vector from each region proposal.




Arbitrary rectangles?
A fixed size input? 227 x 227


5 conv layers + 2 fully connected layers

R-CNN: REGIONS WITH CNN FEATURES


- First Stage: generate category-independent region proposals.
 - 2000 Region proposals every image.
- Second Stage: Extract a fixed length feature vector from each region.
 - a 4096-dimensional feature vector from each region proposal.
- Third Stage: a set of class-specific linear SVMs.
 - object category and localization.



WHY DO WE NEED FAST RCNN?




FAST R-CNN



Idea: No need to recompute features for each box independently.

FAST R-CNN



WHY DO WE NEED FASTER RCNN?

RCNN

- Simple and scalable.
- improves mAP.
- A multistage pipeline.
- Training is expensive in space and time (features are extracted from each region proposal in each image and written into disk).
- Object detection is slow.

Fast-RCNN


- Higher mAP.
- Single stage, end-to-end training.
- No disk storage is required for feature caching.
- proposals are the computational bottleneck in detection systems.

Faster-RCNN

?

FASTER R-CNN



Idea: Integrate the bounding box proposals as part of the CNN predictions.



FASTER R-CNN

Faster-RCNN

Region Proposal Networks:



THEN WHY DOES YOLO EXIST?

RCNN

- Simple and scalable.
- improves mAP.
- A multistage pipeline.
- Training is expensive in space and time (features are extracted from each region proposal in each image and written into disk).
- Object detection is slow.

Fast-RCNN

- Higher mAP.
- Single stage, end-to-end training.
- No disk storage is required for feature caching.
- proposals are the computational bottleneck in detection systems.

Faster-RCNN


- compute proposals with a deep convolutional neural network --*Region Proposal Network (RPN)*
- merge RPN and Fast R-CNN into a single network, enabling nearly cost-free region proposals.

?

YOLO (YOU ONLY LOOK ONCE)

Idea: No bounding box proposal. A single regression problem, straight from image pixels to bounding box coordinates and class probabilities.

- Extremely fast
- Learns generalizable representations.



HANDS ON / CHALLENGE TASK

Task 1: In this task, you will use the provided Fast RCNN and Faster RCNN models to perform object detection on any object detection dataset of your choice. You are expected to fit the dataloaders, train both models, and evaluate their performance.

Dataset: any object detection dataset (Pascal VOC subset, COCO subset, etc)
Metrics: use mean average precision (mAP), inference time, and accuracy.

Reference code for Fast RCNN and Faster RCNN:

https://colab.research.google.com/drive/1E18GeirhjNQrVb_04en5ot_ch4ziH259?usp=sharing